

Evidence before assumptions.

SYNTHETIC SAMPLE - INVENTED DATA - NOT CUSTOMER RESULTS

What a buyer receives

An agreed-scope review of exported temperature, power and reported-throttling telemetry, supported findings, prioritized checks, and a PDF with explicit assumptions. A findings walkthrough is included; format and timing are agreed before payment.

Example scope and data quality

Invented export: 2 GPUs, 10-minute observation window, 1-second sampling, 600 expected rows per GPU. GPU-A has 600 valid rows; GPU-B has 540, leaving 10% missing coverage. Model-specific sensor limits and workload context have not been supplied.

Illustrative findings

01 / Reported thermal throttle

GPU-A reports a thermal-throttle flag in 90 of 600 samples (15%). Maximum reported core temperature: 84 C. This supports investigating a thermal constraint during that window; it does not establish the device safety limit or throughput loss.

02 / Incomplete comparison

GPU-B has a maximum reported core temperature of 72 C and no thermal flags in its 540 recorded samples. Missing intervals prevent a full-window health conclusion.

03 / Power is context, not savings

GPU-A mean reported power: 610 W; GPU-B: 500 W. Equal workload, power-limit configuration and sensor semantics are unverified. The difference does not establish wasted energy or recoverable capacity.

No ROI claim can be made from this sample alone.

A practical next-check plan.

SYNTHETIC SAMPLE - INVENTED DATA - NOT CUSTOMER RESULTS

Priority 1 / Validate the evidence

Confirm the GPU model, sensor units, throttle-reason semantics, timestamps and sampling interval. Check whether flags are instantaneous readings or accumulated counters. Preserve missing intervals rather than silently interpolating them.

Priority 2 / Correlate, then investigate

Align the GPU-A flagged intervals with workload throughput, clock readings, configured power limits and available cooling telemetry. Check whether flags recur under comparable workload conditions. Temperature alone does not identify the cooling fault.

Priority 3 / Re-measure

After the operator approves any intervention, compare matched workload runs using throughput and energy per unit of work. The audit itself makes no GPU, scheduler or cooling-system changes.

Calculations and limitations

Flagged sample share = $90 / 600 = 15\%$. Missing share for GPU-B = $60 / 600 = 10\%$. These are sample shares, not proven durations. No monthly extrapolation, dollar-loss estimate, safety certificate or guaranteed savings is justified by this invented ten-minute export.

How to request the \$80 audit

Email AUDIT to vivaan.thermacompute@gmail.com with GPU model, GPU count and CSV column headers. Do not send credentials, model weights or private logs in the initial request. We confirm suitability, scope, delivery timing and payment instructions before payment.

CrashLens stays free. It organizes supported saved-log signatures and next checks. This separate audit reviews thermal telemetry. Any further service is discussed separately; this offer creates no recurring subscription.